

# **Assignment 2 Report**

110062229 翁語辰

110081014 程詩柔

110062171 陳彥成

## 1. Implementation

- org.vanilladb.bench.server.procedure.as2.UpdatePriceProc  
executeSql(): 把update需要的data(item id, price)用getter從UpdatePriceProcParamHelper取得。然後用StoredProcedureHelper的executeQuery()、和executeUpdate()寫上與JDBC同樣步驟的sql。  
如果沒有先實作這部分, 會出現”no operation implement”的warning, 而且結果跑不出來。
- org.vanilladb.bench.benchmarks.as2.rte.jdbc.As2BenchmarkRte  
getNextTxType(): 利用random的結果和READ\_WRITE\_TX\_TYPE比較來決定下一個txn是update還是read

```
protected As2BenchTransactionType getNextTxType() {  
    double readItemRatio = VanillaBenchParameters.READ_WRITE_TX_RATE;  
  
    if (Math.random() < readItemRatio) {  
        return As2BenchTransactionType.READ_ITEM;  
    } else {  
        return As2BenchTransactionType.UPDATE_PRICE;  
    }  
}
```

getTxExecutor(): 新增一個case, 當type是UPDATE\_PRICE的時候回傳的Executor args裡要放入UpdatePriceParamGen這個class

```
protected As2BenchmarkTxExecutor getTxExecutor(As2BenchTransactionType type) {  
    switch (type) {  
        case READ_ITEM:  
            return new As2BenchmarkTxExecutor(new As2ReadItemParamGen());  
        case UPDATE_PRICE:  
            return new As2BenchmarkTxExecutor(new UpdatePriceParamGen());  
        default:  
            throw new IllegalArgumentException("Unsupported transaction type: " + type);  
    }  
}
```

- org.vanilladb.bench.benchmarks.as2.rte.UpdatePriceParamGen  
generateParameter(): 在arraylist中先依序放入要update的數量(10)、(要update的index、要raise的value) \* 10次, 最後回傳這個array

```
@Override  
public Object[] generateParameter() {  
    RandomValueGenerator rvg = new RandomValueGenerator();  
    ArrayList<Object> paramList = new ArrayList<Object>();  
  
    // Set update count  
    paramList.add(TOTAL_UPDATE_COUNT);  
  
    // Generate item IDs and price raises  
    for (int i = 0; i < TOTAL_UPDATE_COUNT; i++) {  
        // Generate a random item ID to update  
        paramList.add(rvg.number(min:1, As2BenchConstants.NUM_ITEMS));  
  
        // Generate a random price raise value between 0.0 and 5.0  
        paramList.add(rvg.fixedDecimalNumber(decimal:1, min:0.0, max:5.0));  
    }  
  
    return paramList.toArray(new Object[0]);  
}
```

- org.vanilladb.bench.benchmarks.as2.rte.jdbc.As2BenchJdbcExecutor  
execute(): 增加一個case, 當type是UPDATE\_PRICE的時候要呼叫UpdatePriceJdbcJob.execute()

```
@Override  
public SutResultSet execute(Connection conn, As2BenchTransactionType txType, Object[] pars)  
    throws SQLException {  
    switch (txType) {  
        case UPDATE_PRICE:  
            return new UpdatePriceJdbcJob().execute(conn, pars);  
        case TESTBED_LOADER:  
            return new LoadingTestbedJdbcJob().execute(conn, pars);  
        case CHECK_DATABASE:  
            return new CheckDatabaseJdbcJob().execute(conn, pars);  
        case READ_ITEM:  
            return new ReadItemTxnJdbcJob().execute(conn, pars);  
        default:  
            throw new UnsupportedOperationException(  
                String.format("no JDBC implementation for '%s'", txType));  
    }  
}
```

- org.vanilladb.bench.server.param.as2.UpdatePriceProcParamHelper  
大致上和ReadItemProcParamHelper相同, 只是prepareParameters會根據UpdatePriceParamGen產生出來的param格式不同而做不同處理, getResultSetSchema和newResultSetRecord也會因為schema不同而需要做一些更改

```
@Override
public void prepareParameters(Object... pars) {
    int indexCnt = 0;

    updateCount = (Integer) pars[indexCnt++];
    updateItemId = new int[updateCount];
    priceRaise = new double[updateCount];

    for (int i = 0; i < updateCount; i++) {
        updateItemId[i] = (Integer) pars[indexCnt++];
        priceRaise[i] = (Double) pars[indexCnt++];
    }
}

@Override
public Schema getResultSetSchema() {
    Schema sch = new Schema();
    sch.addField(fldName:"update_count", Type.INTEGER);
    return sch;
}

@Override
public SpResultSetRecord newResultSetRecord() {
    SpResultSetRecord rec = new SpResultSetRecord();
    rec.setVal(fldName:"update_count", new IntegerConstant(updateCount));
    return rec;
}
```

- org.vanilladb.bench.benchmarks.as2.rte.jdbc.UpdatePriceJdbcJob  
execute():
  - 用Update price的param helper處理之前param gen產生的param
  - create statement
  - 利用param helper取得每次要更新的i\_id和raise value
  - 和db要目前的i\_price並算出update後的price為何
  - 將該筆資料寫回db
  - 更新完所有item後commit此次transaction

## 2. Screenshot of CSV Report

```
1 time(sec), throughput(txs), avg_latency(ms), min(ms), max(ms), 25th_lat(ms), median_lat(ms), 75th_lat(ms)
2 65,164,12,5,31,9,11,13
3 70,345,11,5,39,9,11,12
4 75,505,11,5,39,8,10,13
5 80,664,12,5,39,8,10,13
6 85,794,12,5,39,8,10,14
7 90,885,13,5,39,8,11,16
8 95,1036,13,5,39,8,11,16
9 100,1228,13,5,51,8,10,15
10 105,1422,12,5,51,8,10,14
11 110,1582,12,5,51,8,10,14
12 115,1696,12,5,51,8,10,15
13 120,1848,12,5,60,8,10,15
```

## 3. Experiments

- Environment:
  - 11th Gen Intel(R) Core(TM) i5-1135G7@ 2.40GHz,16GB RAM,512GB NVMe
  - INTEL,Windows11
- Performance Comparison
  - 90%READ\_ITEM 10%UPDATE\_PRICE

jdbc							
time(sec)	throughput	avg_laten	min(ms)	max(ms)	25th_lat(r	median_l	75th_lat(n
65	365	5	4	14	4	5	5
70	726	5	4	16	4	5	5
75	1098	5	4	17	4	5	5
80	1464	5	4	17	4	5	5
85	1831	5	4	17	4	5	5
90	2188	5	4	18	4	5	5
95	2550	5	4	18	4	5	5
100	2910	5	4	19	4	5	5
105	3263	5	4	19	4	5	5
110	3599	5	4	19	4	5	5
115	3934	5	4	19	4	5	5
120	4263	5	4	19	4	5	5

sp							
time(sec)	throughput(txs)	avg_latency(ms)	min(ms)	max(ms)	25th_lat(ms)	median_lat(ms)	75th_lat(ms)
65	3922	0	0	6	0	0	0
70	8290	0	0	19	0	0	0
75	12362	0	0	19	0	0	0
80	16974	0	0	19	0	0	0
85	20764	0	0	19	0	0	0
90	24792	0	0	30	0	0	0
95	29440	0	0	30	0	0	0
100	33227	0	0	30	0	0	0
105	36746	0	0	30	0	0	0
110	40458	0	0	30	0	0	0
115	43615	0	0	30	0	0	0
120	46746	0	0	30	0	0	0

- 50%READ\_ITEM 50%UPDATE\_PRICE

#### jdbc

time(sec)	throughput	avg_laten	min(ms)	max(ms)	25th_lat(r	median_l	75th_lat(ms)
65	312	6	4	49	5	6	6
70	624	6	4	49	5	6	6
75	929	6	4	49	5	6	7
80	1221	6	4	49	5	6	7
85	1506	6	4	49	5	6	7
90	1804	6	4	49	5	6	7
95	2081	6	4	49	5	6	7
100	2373	6	4	49	5	6	7
105	2639	6	4	49	5	6	7
110	2865	6	4	49	5	6	7
115	3083	7	4	49	5	6	7

#### sp

time(sec)	throughput	avg_laten	min(ms)	max(ms)	25th_lat(r	median_l	75th_lat(ms)
65	3357	0	0	16	0	0	0
70	6859	0	0	27	0	0	0
75	10467	0	0	27	0	0	0
80	14090	0	0	27	0	0	0
85	17709	0	0	27	0	0	0
90	21288	0	0	27	0	0	0
95	24358	0	0	27	0	0	0
100	27485	0	0	27	0	0	0
105	30623	0	0	27	0	0	0
110	33869	0	0	27	0	0	0
115	37144	0	0	27	0	0	0
120	40442	0	0	27	0	0	0

- 0%READ\_ITEM 100%UPDATE\_PRICE

#### jdbc

time(sec)	throughput	avg_laten	min(ms)	max(ms)	25th_lat(r	median_l	75th_lat(ms)
65	231	8	6	31	7	8	9
70	460	8	6	31	7	8	9
75	704	8	6	31	7	8	9
80	937	8	6	31	7	8	9
85	1143	8	6	35	7	8	9
90	1336	8	6	35	7	8	9
95	1544	9	6	35	7	8	9
100	1777	8	6	35	7	8	9
105	2032	8	6	35	7	8	9
110	2290	8	6	35	7	8	9
115	2544	8	6	35	7	8	9

#### sp

time(sec)	throughput	avg_laten	min(ms)	max(ms)	25th_lat(r	median_l	75th_lat(ms)
65	2868	0	0	3	0	0	0
70	5709	0	0	4	0	0	0
75	8582	0	0	4	0	0	0
80	11352	0	0	5	0	0	0
85	14223	0	0	7	0	0	0
90	17228	0	0	7	0	0	0
95	20004	0	0	7	0	0	0
100	22811	0	0	7	0	0	0
105	25747	0	0	7	0	0	0
110	28605	0	0	7	0	0	0
115	31545	0	0	7	0	0	0
120	34657	0	0	7	0	0	0

- Analysis and Explanation

- JDBC vs stored procedure

stored procedure的throughput一向都高出jdbc十幾倍，且latency也比jdbc低很多。那是因為jdbc相較之下，多出許多資料處理。

- 需要連線到server

- 在執行executeQuery時，呼叫 JdbcStatement::executeQuery再呼叫 RemoteStatementImpl::executeQuery 取得transcation後才呼叫 VanillaDb::newPlanner 取得 planner 後呼叫 Planner::createQueryPlan再包成JdbcResultSet 傳送回去(store procedure則是直接呼叫 VanillaDb::newPlanner 取得 planner 後呼叫 Planner::createQueryPlan, 再呼叫 Plan::openVanillaDb::newPlanner 直接完成)

- READ\_WRITE\_TX\_RATE

- 當UPDATE\_PRICE比例從10%提高到100%，throughput會逐漸降低，avg\_latency略為提高，原因是因為UPDATE\_PRICE比READ\_ITEM多做了更新，導致執行時間延長，throughput下降。